

▼ CNN

▼ Import MNIST Images - Deep Learning with PyTorch 14

```
1 import torch
2 import torch.nn as nn
3 import torch.nn.functional as F
4 from torch.utils.data import DataLoader
5 from torchvision import datasets, transforms
6 from torchvision.utils import make_grid
7
8 import numpy as np
9 import pandas as pd
10 from sklearn.metrics import confusion_matrix
11 import matplotlib.pyplot as plt
12
13 %matplotlib inline

1 # Convert MNIST Image Files into a Tensor of 4-Dimensions (# of images, Height, Width, Color Channel)
2 transform = transforms.ToTensor()

1 # Train Data
2 train_data = datasets.MNIST(root='/cnn_data', train=True, download=True, transform=transform)

1 # test Data
2 test_data = datasets.MNIST(root='/cnn_data', train=False, download=True, transform=transform)

1 train_data

Dataset MNIST
  Number of datapoints: 60000
  Root location: /cnn_data
  Split: Train
  StandardTransform
  Transform: ToTensor()

1 test_data

Dataset MNIST
  Number of datapoints: 10000
  Root location: /cnn_data
  Split: Test
  StandardTransform
  Transform: ToTensor()

1 pwd

'/content'

1 ls

sample_data/
```

```
1 cd ../
```

```
/
```

```
1 ls -al
```

```
total 116
drwxr-xr-x 1 root root 4096 Oct 17 19:12 ./
drwxr-xr-x 1 root root 4096 Oct 17 19:12 ../
lrwxrwxrwx 1 root root    7 Jun  5 14:02 bin -> usr/bin/
drwxr-xr-x 2 root root 4096 Apr 18  2022 boot/
drwxr-xr-x 3 root root 4096 Oct 17 19:12 cnn_data/
drwxr-xr-x 1 root root 4096 Oct 16 13:23 content/
-rw-r--r-- 1 root root 4332 Jun 21 00:40 cuda-keyring_1.0-1_all.deb
drwxr-xr-x 1 root root 4096 Oct 16 13:52 datalab/
drwxr-xr-x 5 root root  360 Oct 17 19:11 dev/
-rwxr-xr-x 1 root root    0 Oct 17 19:11 .dockerenv*
drwxr-xr-x 1 root root 4096 Oct 17 19:11 etc/
drwxr-xr-x 2 root root 4096 Apr 18  2022 home/
lrwxrwxrwx 1 root root    7 Jun  5 14:02 lib -> usr/lib/
lrwxrwxrwx 1 root root    9 Jun  5 14:02 lib32 -> usr/lib32/
lrwxrwxrwx 1 root root    9 Jun  5 14:02 lib64 -> usr/lib64/
lrwxrwxrwx 1 root root   10 Jun  5 14:02 libx32 -> usr/libx32/
drwxr-xr-x 2 root root 4096 Jun  5 14:02 media/
drwxr-xr-x 2 root root 4096 Jun  5 14:02 mnt/
-rw-r--r-- 1 root root 17294 Jun 21 00:39 NGC-DL-CONTAINER-LICENSE
drwxr-xr-x 1 root root 4096 Oct 16 13:53 opt/
dr-xr-xr-x 175 root root    0 Oct 17 19:11 proc/
drwxr-xr-x 15 root root 4096 Oct 16 13:20 python-apt/
drwx----- 1 root root 4096 Oct 16 13:53 root/
drwxr-xr-x 1 root root 4096 Oct 16 13:15 run/
lrwxrwxrwx 1 root root    8 Jun  5 14:02 sbin -> usr/sbin/
drwxr-xr-x 2 root root 4096 Jun  5 14:02 srv/
dr-xr-xr-x 13 root root    0 Oct 17 19:11 sys/
drwxrwxrwt 1 root root 4096 Oct 17 19:11 tmp/
drwxr-xr-x 1 root root 4096 Oct 16 13:39 tools/
drwxr-xr-x 1 root root 4096 Oct 16 13:53 usr/
drwxr-xr-x 1 root root 4096 Oct 16 13:52 var/
```

```
1 cd cnn_data
```

```
/cnn_data
```

```
1 ls -l
```

```
total 4
drwxr-xr-x 3 root root 4096 Oct 17 19:12 MNIST/
```

```
1 cd /
```

```
/
```

```
1 ls -l
```

```
total 108
lrwxrwxrwx 1 root root    7 Jun  5 14:02 bin -> usr/bin/
drwxr-xr-x 2 root root 4096 Apr 18  2022 boot/
drwxr-xr-x 3 root root 4096 Oct 17 19:12 cnn_data/
drwxr-xr-x 1 root root 4096 Oct 16 13:23 content/
-rw-r--r-- 1 root root 4332 Jun 21 00:40 cuda-keyring_1.0-1_all.deb
drwxr-xr-x 1 root root 4096 Oct 16 13:52 datalab/
```

```

drwxr-xr-x  5 root root   360 Oct 17 19:11 dev/
drwxr-xr-x  1 root root  4096 Oct 17 19:11 etc/
drwxr-xr-x  2 root root  4096 Apr 18  2022 home/
lrwxrwxrwx  1 root root    7 Jun  5 14:02 lib -> usr/lib/
lrwxrwxrwx  1 root root    9 Jun  5 14:02 lib32 -> usr/lib32/
lrwxrwxrwx  1 root root    9 Jun  5 14:02 lib64 -> usr/lib64/
lrwxrwxrwx  1 root root   10 Jun  5 14:02 libx32 -> usr/libx32/
drwxr-xr-x  2 root root  4096 Jun  5 14:02 media/
drwxr-xr-x  2 root root  4096 Jun  5 14:02 mnt/
-rw-r--r--  1 root root 17294 Jun 21 00:39 NGC-DL-CONTAINER-LICENSE
drwxr-xr-x  1 root root  4096 Oct 16 13:53 opt/
dr-xr-xr-x 175 root root    0 Oct 17 19:11 proc/
drwxr-xr-x 15 root root  4096 Oct 16 13:20 python-apt/
drwx----- 1 root root  4096 Oct 16 13:53 root/
drwxr-xr-x  1 root root  4096 Oct 16 13:15 run/
lrwxrwxrwx  1 root root    8 Jun  5 14:02/sbin -> usr/sbin/
drwxr-xr-x  2 root root  4096 Jun  5 14:02 srv/
dr-xr-xr-x 13 root root    0 Oct 17 19:11 sys/
drwxrwxrwt  1 root root  4096 Oct 17 19:11 tmp/
drwxr-xr-x  1 root root  4096 Oct 16 13:39 tools/
drwxr-xr-x  1 root root  4096 Oct 16 13:53 usr/
drwxr-xr-x  1 root root  4096 Oct 16 13:52 var/

```

```
1 cd content/
```

```
    /content
```

```
1 ls -al
```

```

total 16
drwxr-xr-x 1 root root 4096 Oct 16 13:23 ./
drwxr-xr-x 1 root root 4096 Oct 17 19:12 ../
drwxr-xr-x 4 root root 4096 Oct 16 13:23 .config/
drwxr-xr-x 1 root root 4096 Oct 16 13:23 sample_data/

```

▼ Convolutional and Pooling Layers - Deep Learning with PyTorch 15

Double-cliquez (ou appuyez sur Entrée) pour modifier

```

1 # Create a small batch size for images... let's say 10
2 train_loader = DataLoader(train_data, batch_size=10, shuffle=True)
3 test_loader = DataLoader(test_data, batch_size=10, shuffle=False)

1 # Define the CNN Model
2 # Describe the convolutional layer and what it's doing (2 convolutional layers)
3 # This is an example
4 conv1 = nn.Conv2d(1, 6, 3, 1)
5 conv2 = nn.Conv2d(in_channels=6, out_channels=16, kernel_size=3, stride=1)
6

1 # Grab 1 MNIST record/image
2 for i, (X_train, y_train) in enumerate(train_data):
3     break

1 X_train

tensor([[[[0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000,

```

```

0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000,
0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000,
0.0000, 0.0000, 0.0000, 0.0000],
[0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000,
0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000,
0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000,
0.0000, 0.0000, 0.0000, 0.0000],
[0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000,
0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000,
0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000,
0.0000, 0.0000, 0.0000, 0.0000],
[0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000,
0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000,
0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000,
0.0000, 0.0000, 0.0000, 0.0000],
[0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000,
0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000,
0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000,
0.0000, 0.0000, 0.0000, 0.0000],
[0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000,
0.0000, 0.0000, 0.0000, 0.0000, 0.0118, 0.0706, 0.0706, 0.0706,
0.4941, 0.5333, 0.6863, 0.1020, 0.6510, 1.0000, 0.9686, 0.4980,
0.0000, 0.0000, 0.0000, 0.0000],
[0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000,
0.1176, 0.1412, 0.3686, 0.6039, 0.6667, 0.9922, 0.9922, 0.9922,
0.9922, 0.9922, 0.8824, 0.6745, 0.9922, 0.9490, 0.7647, 0.2510,
0.0000, 0.0000, 0.0000, 0.0000],
[0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.1922,
0.9333, 0.9922, 0.9922, 0.9922, 0.9922, 0.9922, 0.9922, 0.9922,
0.9922, 0.9843, 0.3647, 0.3216, 0.3216, 0.2196, 0.1529, 0.0000,
0.0000, 0.0000, 0.0000, 0.0000],
[0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0706,
0.8588, 0.9922, 0.9922, 0.9922, 0.9922, 0.9922, 0.7765, 0.7137,
0.9686, 0.9451, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000,
0.0000, 0.0000, 0.0000, 0.0000],
[0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000,
0.3137, 0.6118, 0.4196, 0.9922, 0.9922, 0.8039, 0.0431, 0.0000,
0.1686, 0.6039, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000,
0.0000, 0.0000, 0.0000, 0.0000],
[0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000,
0.0000, 0.0549, 0.0039, 0.6039, 0.9922, 0.3529, 0.0000, 0.0000,
0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000,
0.0000, 0.0000, 0.0000, 0.0000],
[0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000,
0.0000, 0.0000, 0.0000, 0.5451, 0.9922, 0.7451, 0.0078, 0.0000,
0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000,
0.0000, 0.0000, 0.0000, 0.0000],
[0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000,
0.0000, 0.0000, 0.0000, 0.0431, 0.7451, 0.9922, 0.2745, 0.0000,
0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000,
0.0000, 0.0000, 0.0000, 0.0000],
[0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000,
0.0000, 0.0000, 0.0000, 0.0000, 0.1373, 0.9451, 0.8824, 0.6275,
0.4235, 0.0039, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000,
0.0000, 0.0000, 0.0000, 0.0000],
[0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000,
0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000,
0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.3176, 0.9412, 0.9922,

```

```
1 X_train.shape
```

```
torch.Size([1, 28, 28])
```

```
1 x = X_train.view(1,1, 28, 28)
```

```

1 # Perform the first convolution
2 x = F.relu(conv1(x)) # Rectified Linear Unit for the activation function

1 x

tensor([[[[0.1709, 0.1709, 0.1709, ..., 0.1709, 0.1709, 0.1709],
          [0.1709, 0.1709, 0.1709, ..., 0.1709, 0.1709, 0.1709],
          [0.1709, 0.1709, 0.1709, ..., 0.1709, 0.1709, 0.1709],
          ...,
          [0.1709, 0.1709, 0.1158, ..., 0.1709, 0.1709, 0.1709],
          [0.1709, 0.1709, 0.1537, ..., 0.1709, 0.1709, 0.1709],
          [0.1709, 0.1709, 0.1709, ..., 0.1709, 0.1709, 0.1709]],

        [[0.0006, 0.0006, 0.0006, ..., 0.0006, 0.0006, 0.0006],
          [0.0006, 0.0006, 0.0006, ..., 0.0006, 0.0006, 0.0006],
          [0.0006, 0.0006, 0.0006, ..., 0.0006, 0.0006, 0.0006],
          ...,
          [0.0006, 0.0006, 0.1298, ..., 0.0006, 0.0006, 0.0006],
          [0.0006, 0.0006, 0.0454, ..., 0.0006, 0.0006, 0.0006],
          [0.0006, 0.0006, 0.0006, ..., 0.0006, 0.0006, 0.0006]],

        [[0.1699, 0.1699, 0.1699, ..., 0.1699, 0.1699, 0.1699],
          [0.1699, 0.1699, 0.1699, ..., 0.1699, 0.1699, 0.1699],
          [0.1699, 0.1699, 0.1699, ..., 0.1699, 0.1699, 0.1699],
          ...,
          [0.1699, 0.1699, 0.3766, ..., 0.1699, 0.1699, 0.1699],
          [0.1699, 0.1699, 0.3260, ..., 0.1699, 0.1699, 0.1699],
          [0.1699, 0.1699, 0.1699, ..., 0.1699, 0.1699, 0.1699]],

        [[0.0000, 0.0000, 0.0000, ..., 0.0000, 0.0000, 0.0000],
          [0.0000, 0.0000, 0.0000, ..., 0.0000, 0.0000, 0.0000],
          [0.0000, 0.0000, 0.0000, ..., 0.0000, 0.0000, 0.0000],
          ...,
          [0.0000, 0.0000, 0.0000, ..., 0.0000, 0.0000, 0.0000],
          [0.0000, 0.0000, 0.0000, ..., 0.0000, 0.0000, 0.0000],
          [0.0000, 0.0000, 0.0000, ..., 0.0000, 0.0000, 0.0000]],

        [[0.0591, 0.0591, 0.0591, ..., 0.0591, 0.0591, 0.0591],
          [0.0591, 0.0591, 0.0591, ..., 0.0591, 0.0591, 0.0591],
          [0.0591, 0.0591, 0.0591, ..., 0.0591, 0.0591, 0.0591],
          ...,
          [0.0591, 0.0591, 0.0000, ..., 0.0591, 0.0591, 0.0591],
          [0.0591, 0.0591, 0.1010, ..., 0.0591, 0.0591, 0.0591],
          [0.0591, 0.0591, 0.0591, ..., 0.0591, 0.0591, 0.0591]],

        [[0.3034, 0.3034, 0.3034, ..., 0.3034, 0.3034, 0.3034],
          [0.3034, 0.3034, 0.3034, ..., 0.3034, 0.3034, 0.3034],
          [0.3034, 0.3034, 0.3034, ..., 0.3034, 0.3034, 0.3034],
          ...,
          [0.3034, 0.3034, 0.4017, ..., 0.3034, 0.3034, 0.3034],
          [0.3034, 0.3034, 0.4095, ..., 0.3034, 0.3034, 0.3034],
          [0.3034, 0.3034, 0.3034, ..., 0.3034, 0.3034, 0.3034]]]],
       grad_fn=<ReluBackward0>)

1 # 1 is the single image, 6 is the filters asked for, 26x26
2 x.shape

torch.Size([1, 6, 26, 26])

1 # Pass through the pooling layer
2 x = F.max_pool2d(x, 2, 2) # Kernel of 2 and stride of 2

```

```

1 x.shape # 26 / 2 = 13

    torch.Size([1, 6, 13, 13])

1 # Do the second convolutional layer
2 x = F.relu(conv2(x))

1 x.shape # Again, no padding was specified so 2 pixels were lost around the outside of the image

    torch.Size([1, 16, 11, 11])

1 # Pooling layer
2 x = F.max_pool2d(x, 2, 2)

1 x.shape # 11 / 2 = 5.5 but it is rounded down because no data can be invented to round up

    torch.Size([1, 16, 5, 5])

1 (((28-2) / 2) - 2) / 2

    5.5

```

▼ Convolutional Neural Network Model - Deep Learning with PyTorch 16

```

1 # Model Class
2 class ConvolutionalNetwork(nn.Module):
3     def __init__(self) -> None:
4         super().__init__()
5         self.conv1 = nn.Conv2d(1, 6, 3, 1)
6         self.conv2 = nn.Conv2d(6, 16, 3, 1)
7
8         # Fully Connected Layers
9         self.fc1 = nn.Linear(5*5*16, 120)
10        self.fc2 = nn.Linear(120, 84)
11        self.fc3 = nn.Linear(84, 10)
12
13    def forward(self, X):
14        X = F.relu(self.conv1(X))
15        X = F.max_pool2d(X, 2, 2) # 2x2 kernel and stride = 2
16        # Second pass
17        X = F.relu(self.conv2(X))
18        X = F.max_pool2d(X, 2, 2) # 2x2 kernel and stride = 2
19
20        # Re-View the data to flatten it out
21        X = X.view(-1, 16*5*5) # Negative one so the batch size can be varied
22
23        # Fully Connected Layers
24        X = F.relu(self.fc1(X))
25        X = F.relu(self.fc2(X))
26        X = self.fc3(X)
27
28        return F.log_softmax(X, dim=1)
29
30

1 # Create an Instance of the Model
2 torch.manual_seed(41)
3 model = ConvolutionalNetwork()

```

4 model

```
ConvolutionalNetwork(
  (conv1): Conv2d(1, 6, kernel_size=(3, 3), stride=(1, 1))
  (conv2): Conv2d(6, 16, kernel_size=(3, 3), stride=(1, 1))
  (fc1): Linear(in_features=400, out_features=120, bias=True)
  (fc2): Linear(in_features=120, out_features=84, bias=True)
  (fc3): Linear(in_features=84, out_features=10, bias=True)
)
```

```
1 # Loss Function Optimizer
2 criterion = nn.CrossEntropyLoss()
3 optimizer = torch.optim.Adam(model.parameters(), lr=0.001) # The smaller the learning rate, the longer it's g
```

▼ Train and Test CNN Model - Deep Learning with PyTorch 17

```
1 import time
2 start_time = time.time()
3
4 # Create Variables to track things
5 epochs = 5
6 train_losses = []
7 test_losses = []
8 train_correct = []
9 test_correct = []
10
11
12 # For Loop offor Epochs
13 for i in range(epochs):
14     training_correct = 0
15     testing_correct = 0
16
17     # Train
18     for b, (X_train, y_train) in enumerate(train_loader):
19         b += 1 # Start the batches at 1
20
21         y_pred = model(X_train) # Get the predicted values from the training set (data is 2d, not flattened.)
22         loss = criterion(y_pred, y_train) # How off are we? Compare the predictions to the correct answers in y_1
23
24         predicted = torch.max(y_pred.data, 1)[1] # Add up the number of correct predictions. Indexed off the first
25         batch_correct = (predicted == y_train).sum() # How many we got correct from this specific batch. True=1,
26         training_correct += batch_correct # Keep track as we go along in training.
27
28         # Update the parameters
29         optimizer.zero_grad()
30         loss.backward()
31         optimizer.step()
32
33         # Print out some results
34         if b % 600 == 0 :
35             print(f'Epoch: {i} Batch: {b} Loss: {loss.item()}')
36
37     train_losses.append(loss)
38     train_correct.append(training_correct)
39
40
41     # Test
42     with torch.no_grad(): # No gradient so the weights and the bias are not updated with test data
43         for b, (X_test, y_test) in enumerate(test_loader):
44             y_val = model(X_test)
45             predicted = torch.max(y_val.data, 1)[1] # Adding up correct predictions
46             testing_correct += (predicted == y_test).sum() # True=1, False=0, sum all
47
```

```
48     loss = criterion(y_val, y_test)
49     test_losses.append(loss)
50     test_correct.append(testing_correct)
51
52     current_time = time.time()
53     total = current_time - start_time
54     print(f'Training time: {total/60} minutes!')
55     total

Epoch: 0 Batch: 600 Loss: 4.9471535021439195e-05
Epoch: 0 Batch: 1200 Loss: 1.8047001503873616e-05
Epoch: 0 Batch: 1800 Loss: 8.344646573732462e-08
Epoch: 0 Batch: 2400 Loss: 0.003162816632539034
Epoch: 0 Batch: 3000 Loss: 1.1622306374192704e-05
Epoch: 0 Batch: 3600 Loss: 0.0001164354252978228
Epoch: 0 Batch: 4200 Loss: 9.417489081897656e-07
Epoch: 0 Batch: 4800 Loss: 0.00047675552195869386
Epoch: 0 Batch: 5400 Loss: 0.0
Epoch: 0 Batch: 6000 Loss: 1.609314722372801e-06
Epoch: 1 Batch: 600 Loss: 2.777537929432583e-06
Epoch: 1 Batch: 1200 Loss: 1.5139510196604533e-06
Epoch: 1 Batch: 1800 Loss: 2.731903805397451e-05
Epoch: 1 Batch: 2400 Loss: 1.889325176307466e-05
Epoch: 1 Batch: 3000 Loss: 1.4351843674376141e-05
Epoch: 1 Batch: 3600 Loss: 9.142941962636542e-06
Epoch: 1 Batch: 4200 Loss: 8.943313878262416e-05
Epoch: 1 Batch: 4800 Loss: 0.00012190106644993648
Epoch: 1 Batch: 5400 Loss: 1.3207888514443766e-05
Epoch: 1 Batch: 6000 Loss: 0.0
Epoch: 2 Batch: 600 Loss: 2.1097897842992097e-05
Epoch: 2 Batch: 1200 Loss: 0.0
Epoch: 2 Batch: 1800 Loss: 8.022654583328404e-06
Epoch: 2 Batch: 2400 Loss: 0.0
Epoch: 2 Batch: 3000 Loss: 0.0
Epoch: 2 Batch: 3600 Loss: 0.0013853703858330846
Epoch: 2 Batch: 4200 Loss: 8.702245395397767e-07
Epoch: 2 Batch: 4800 Loss: 0.0
Epoch: 2 Batch: 5400 Loss: 5.555012648983393e-06
Epoch: 2 Batch: 6000 Loss: 3.576274423267023e-07
Epoch: 3 Batch: 600 Loss: 9.775114904186921e-07
Epoch: 3 Batch: 1200 Loss: 1.1514954167068936e-05
Epoch: 3 Batch: 1800 Loss: 7.748575399091351e-07
Epoch: 3 Batch: 2400 Loss: 2.2887920749781188e-06
Epoch: 3 Batch: 3000 Loss: 1.192092824453539e-08
Epoch: 3 Batch: 3600 Loss: 0.0
Epoch: 3 Batch: 4200 Loss: 0.002368538174778223
Epoch: 3 Batch: 4800 Loss: 0.0001560908422106877
Epoch: 3 Batch: 5400 Loss: 4.863749927608296e-05
Epoch: 3 Batch: 6000 Loss: 1.5114685993466992e-05
Epoch: 4 Batch: 600 Loss: 0.0
Epoch: 4 Batch: 1200 Loss: 5.960462701182223e-08
Epoch: 4 Batch: 1800 Loss: 2.384185648907078e-08
Epoch: 4 Batch: 2400 Loss: 1.2159273410361493e-06
Epoch: 4 Batch: 3000 Loss: 0.0002526980242691934
Epoch: 4 Batch: 3600 Loss: 0.0036067436449229717
Epoch: 4 Batch: 4200 Loss: 0.0
Epoch: 4 Batch: 4800 Loss: 3.6161560274194926e-05
Epoch: 4 Batch: 5400 Loss: 0.4160960614681244
Epoch: 4 Batch: 6000 Loss: 4.497986446949653e-05
Training time: 4.19429939587911 minutes!
251.65796375274658
```

```
1
2 import time
3 start_time = time.time()
```



```

4
5 # Create Variables To Tracks Things
6 epochs = 5
7 train_losses = []
8 test_losses = []
9 train_correct = []
10 test_correct = []
11
12 # For Loop of Epochs
13 for i in range(epochs):
14     trn_corr = 0
15     tst_corr = 0
16
17
18 # Train
19 for b,(X_train, y_train) in enumerate(train_loader):
20     b+=1 # start our batches at 1
21     y_pred = model(X_train) # get predicted values from the training set. Not flattened 2D
22     loss = criterion(y_pred, y_train) # how off are we? Compare the predictions to correct answers in y_train
23
24     predicted = torch.max(y_pred.data, 1)[1] # add up the number of correct predictions. Indexed off the first
25     batch_corr = (predicted == y_train).sum() # how many we got correct from this batch. True = 1, False=0, s
26     trn_corr += batch_corr # keep track as we go along in training.
27
28     # Update our parameters
29     optimizer.zero_grad()
30     loss.backward()
31     optimizer.step()
32
33
34     # Print out some results
35     if b%600 == 0:
36         print(f'Epoch: {i} Batch: {b} Loss: {loss.item()}')
37
38     train_losses.append(loss)
39     train_correct.append(trn_corr)
40
41
42 # Test
43 with torch.no_grad(): #No gradient so we don't update our weights and biases with test data
44     for b,(X_test, y_test) in enumerate(test_loader):
45         y_val = model(X_test)
46         predicted = torch.max(y_val.data, 1)[1] # Adding up correct predictions
47         tst_corr += (predicted == y_test).sum() # T=1 F=0 and sum away
48
49
50     loss = criterion(y_val, y_test)
51     test_losses.append(loss)
52     test_correct.append(tst_corr)
53
54
55
56 current_time = time.time()
57 total = current_time - start_time
58 print(f'Training Took: {total/60} minutes!')

```

Epoch: 0 Batch: 600 Loss: 0.0027929155621677637
Epoch: 0 Batch: 1200 Loss: 3.493723488645628e-05
Epoch: 0 Batch: 1800 Loss: 0.0010003604693338275
Epoch: 0 Batch: 2400 Loss: 0.1295616179704666
Epoch: 0 Batch: 3000 Loss: 1.3505514289136045e-05
Epoch: 0 Batch: 3600 Loss: 0.00021176428708713502
Epoch: 0 Batch: 4200 Loss: 0.00011720164911821485
Epoch: 0 Batch: 4800 Loss: 0.0009881147416308522
Epoch: 0 Batch: 5400 Loss: 2.384185648907078e-08
Epoch: 0 Batch: 6000 Loss: 0.0
Epoch: 1 Batch: 600 Loss: 0.0008038681116886437

```
Epoch: 1 Batch: 1200 Loss: 4.76837058727142e-08
Epoch: 1 Batch: 1800 Loss: 3.4093500289600343e-06
Epoch: 1 Batch: 2400 Loss: 3.4710730687947944e-05
Epoch: 1 Batch: 3000 Loss: 0.0002807883720379323
Epoch: 1 Batch: 3600 Loss: 9.417489081897656e-07
Epoch: 1 Batch: 4200 Loss: 0.008015107363462448
Epoch: 1 Batch: 4800 Loss: 0.00021272152662277222
Epoch: 1 Batch: 5400 Loss: 7.617282335559139e-06
Epoch: 1 Batch: 6000 Loss: 7.709871715633199e-05
Epoch: 2 Batch: 600 Loss: 1.1241128959227353e-05
Epoch: 2 Batch: 1200 Loss: 9.536717584524013e-07
Epoch: 2 Batch: 1800 Loss: 0.0005367971025407314
Epoch: 2 Batch: 2400 Loss: 4.148400421399856e-06
Epoch: 2 Batch: 3000 Loss: 0.0
Epoch: 2 Batch: 3600 Loss: 0.07778234779834747
Epoch: 2 Batch: 4200 Loss: 1.001354917207209e-06
Epoch: 2 Batch: 4800 Loss: 0.042573653161525726
Epoch: 2 Batch: 5400 Loss: 0.0
Epoch: 2 Batch: 6000 Loss: 0.00151331617962569
Epoch: 3 Batch: 600 Loss: 1.4113868928689044e-05
Epoch: 3 Batch: 1200 Loss: 2.2530307433044072e-06
Epoch: 3 Batch: 1800 Loss: 2.384185471271394e-08
Epoch: 3 Batch: 2400 Loss: 3.576272717964457e-07
Epoch: 3 Batch: 3000 Loss: 2.384185648907078e-08
Epoch: 3 Batch: 3600 Loss: 0.054267413914203644
Epoch: 3 Batch: 4200 Loss: 0.680134654045105
Epoch: 3 Batch: 4800 Loss: 0.0008131394279189408
Epoch: 3 Batch: 5400 Loss: 1.3255626072350424e-05
Epoch: 3 Batch: 6000 Loss: 1.0728830091011332e-07
Epoch: 4 Batch: 600 Loss: 0.000250897224759683
Epoch: 4 Batch: 1200 Loss: 0.0004779839946422726
Epoch: 4 Batch: 1800 Loss: 0.0010529584251344204
Epoch: 4 Batch: 2400 Loss: 0.001856612740084529
Epoch: 4 Batch: 3000 Loss: 1.2707006135315169e-05
Epoch: 4 Batch: 3600 Loss: 0.0
Epoch: 4 Batch: 4200 Loss: 1.192092824453539e-08
Epoch: 4 Batch: 4800 Loss: 3.576278118089249e-08
Epoch: 4 Batch: 5400 Loss: 2.0265562739041343e-07
Epoch: 4 Batch: 6000 Loss: 0.0012134775752201676
Training Took: 4.140049517154694 minutes!
```

